

# Workshop 2 - Predicting loan cases using Decision Tree

Student Number: 2302546

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: dataset = pd.read_csv("train_ctrUa4K.csv")
```

```
In [3]: dataset.head()
```

```
Out[3]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

```
In [4]: dataset.shape
```

```
Out[4]: (614, 13)
```

```
In [5]: dataset.ndim
```

```
Out[5]: 2
```

## Generating unique dataset for this task

I will be generating a unique dataset for this notebook using the last two digits of my student number in the random\_state.

```
In [6]: dataset.size
```

```
Out[6]: 7982
```

```
In [7]: dataset = dataset.sample(n=550, random_state = 46)
```

```
In [8]: dataset.to_csv('Arabambi_2302546.csv')
```

```
In [9]: data = pd.read_csv('Arabambi_2302546.csv')
```

In [10]: `data.head()`

Out[10]:

	Unnamed: 0	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome
0	273	LP001894	Male	Yes	0	Graduate	No	2620
1	246	LP001814	Male	Yes	2	Graduate	No	9703
2	371	LP002197	Male	Yes	2	Graduate	No	5185
3	538	LP002739	Male	Yes	0	Not Graduate	No	2917
4	288	LP001931	Female	No	0	Graduate	No	4124

In [11]: `data=data.drop('Unnamed: 0', axis = 1)`

In [12]: `data.head()`

Out[12]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	

## Data Visualization

### Q1. Use and explain the following DataFrame functions/properties on your data.

- `describe()`
- `size`
- `ndim`
- `shape`

In [13]: `print(data.describe())`

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	550.000000	550.000000	531.000000	537.000000
mean	5497.836364	1610.992582	148.527307	341.787709
std	6263.552850	2938.590680	88.188674	64.796521
min	150.000000	0.000000	9.000000	12.000000
25%	2843.000000	0.000000	100.000000	360.000000
50%	3815.000000	1149.000000	128.000000	360.000000
75%	5844.000000	2297.250000	170.000000	360.000000
max	81000.000000	41667.000000	700.000000	480.000000

	Credit_History
count	506.000000
mean	0.839921
std	0.367042
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

The `describe()` function is a pandas DataFrame method used to generate descriptive statistics of a DataFrame. It computes the count, mean, standard deviation, minimum and maximum values, as well as the quartiles (25%, 50%, and 75%) of each numerical column in the DataFrame.

In the given DataFrame, `data.describe()` has generated the summary statistics for all the columns present in the DataFrame, including Gender, Married, Dependents, Education, Self\_Employed, ApplicantIncome, CoapplicantIncome, Loan\_Amount\_Term, Credit\_History, Property\_Area, Loan\_Status, LoanAmount\_log, Loan\_Amount\_Term\_log, and TotalIncome.

```
In [14]: data.size
```

```
Out[14]: 7150
```

'data.size' is a property of the DataFrame object in pandas which returns the number of elements in the DataFrame. This property returns the same value as the product of the number of rows and columns in the DataFrame. In other words, it returns the total number of cells or entries in the DataFrame.

In the code above, 'data.size' returns the total number of elements in the 'data' DataFrame, which is equal to 7150.

```
In [15]: data.ndim
```

```
Out[15]: 2
```

`data.ndim` returns the number of dimensions of the dataframe or axes of the DataFrame. For this pandas dataframe, it returns 2, as it is a 2-dimensional data structure.

```
In [16]: data.shape
```

```
Out[16]: (550, 13)
```

`data.shape` returns the dimensions of the dataframe in the form of a tuple. The first element shows the number of rows(550) and the second element shows the number of columns(13) in the dataframe.

## Q2. Is there any difference between dimensions of the original dataset and the new dataset? If yes, what is the difference?

Yes, there are differences between the dimensions of the old and new datasets(Comparing In[4] and In[6] to In[14] and In[16]). The size and shape of the data set has been altered in In[7]. However, the dimensions (ndim) remains the same.

## Q3. What are the possible values 'Education' can take? Write code to display all the possible values of 'Education'.

```
In [17]: print(data['Education'].unique())  
[ 'Graduate' 'Not Graduate' ]
```

```
In [18]: data['Education'].value_counts()
```

```
Out[18]: Graduate      438  
Not Graduate    112  
Name: Education, dtype: int64
```

The values 'Education' can take is

- Graduate
- Not Graduate

In[17] and In[18] shows this.

## Data Analysis

```
In [19]: columns = data.columns  
columns
```

```
Out[19]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',  
            'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
            'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],  
          dtype='object')
```


```
In [20]: data.head()
```

Out[20]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	

In [21]: `data['ApplicantIncome'].hist(bins=50)`

Out[21]: <AxesSubplot:>



## Q4. Use boxplot and histogram on 'ApplicantIncome' to visualise its distribution.

Histogram and boxplot are used on the same feature to visualise the data distribution. Compare both the plots and report:


4a. What are the extreme values? Are there any outliers(s) exist in this dataset? Explain with example based on the 'ApplicantIncome'?

4b. Are the results of both the plots comparable? Are there any differences in the two plots? What are the key differences?

# Answer


```
In [22]: data.boxplot(column='ApplicantIncome')
```

```
Out[22]: <AxesSubplot:>
```



```
In [23]: data['ApplicantIncome'].hist(bins=50)
```

```
Out[23]: <AxesSubplot:>
```



Q4a. Extreme Values defines the boundaries of a normal distribution, any data that falls outside the maximum and minimum point in the normal distribution is referred to as an outlier. There are several outliers in the dataset all coming in above the maximum value of this dataset at 10000+ as show in Out[22]


Q4b. Both the plots give different information about the distribution of the data. The histogram shows the frequency distribution of the data, while the boxplot shows the summary statistics such as median, first and third quartile, and outliers. The histogram gives a rough idea of the data distribution, while the boxplot gives a more concise and compact representation of the data distribution especially when defining the boundaries and identifying the outliers present in the dataset.

## Try-It-Yourself

### Use Histogram and Box plot on 'LoanAmount' and observe extreme values


```
In [24]: #Box plot  
data.boxplot(column='LoanAmount')
```

```
Out[24]: <AxesSubplot:>
```



```
In [25]: #Histogram  
data['LoanAmount'].hist(bins=50)
```

```
Out[25]: <AxesSubplot:>
```



Extreme values for LoanAmount can be observed using the boxplot and histogram, the extreme values are about 10 and 270. Also, median is about 120 in the box plot. The box plot shows these values way more noticeably than the histogram.

## Categorical Variable Analysis

```
In [26]: data['Credit_History'].value_counts()
```

```
Out[26]: 1.0    425
         0.0    81
         Name: Credit_History, dtype: int64
```

```
In [27]: credit_history = data['Credit_History'].value_counts(ascending=True)

         loan_probability = data.pivot_table(values='Loan_Status', index=['Credit_History'],
         aggfunc=lambda x: x.map({'Y':1, 'N':0}).mean())

         print('Frequency Table for Credit History:')
         print(credit_history)
         print('\nProbability of getting loan for each Credit History class:')
         print(loan_probability)
```

```
Frequency Table for Credit History:
0.0    81
1.0   425
Name: Credit_History, dtype: int64
```

```
Probability of getting loan for each Credit History class:
Loan_Status
Credit_History
0.0           0.086420
1.0           0.790588
```

```
In [28]: data['Loan_Status'].value_counts()
```



```
Out[28]: Y      376
         N      174
         Name: Loan_Status, dtype: int64
```

```
In [29]: data.shape
```

```
Out[29]: (550, 13)
```


```
In [30]: data.head()
```


```
Out[30]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	

```
In [31]: fig = plt.figure(figsize=(8,4))
         ax1 = fig.add_subplot(121)
         ax1.set_xlabel('Credit_History')
         ax1.set_ylabel('Count of Applicants')
         ax1.set_title("Applicants by Credit_History")
         credit_history.plot(kind='bar')
         plt.show()

         ax2 = fig.add_subplot(122)
         ax2.set_xlabel('Credit_History')
         ax2.set_ylabel('Probability of getting loan')
         ax2.set_title("Probability of getting loan by credit history")
         loan_probability.plot(kind = 'bar')
         plt.show()
```





## Data Pre-Processing

```
In [32]: data['Gender'].value_counts()
```

```
Out[32]: Male      437  
Female    100  
Name: Gender, dtype: int64
```

### Filling in missing values by mean

```
In [33]: data.apply(lambda x: sum(x.isnull()), axis=0)
```

```
Out[33]: Loan_ID      0  
Gender      13  
Married      3  
Dependents   13  
Education    0  
Self_Employed 28  
ApplicantIncome 0  
CoapplicantIncome 0  
LoanAmount   19  
Loan_Amount_Term 13  
Credit_History 44  
Property_Area 0  
Loan_Status  0  
dtype: int64
```

```
In [34]: data.head()
```

Out[34]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	

In [35]: `data['LoanAmount'].fillna(data['LoanAmount'].mean(), inplace = True)`

In [36]: `data.head()`

Out[36]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	

In [37]: `data.apply(lambda x: sum(x.isnull()), axis=0)`

Out[37]:

Loan_ID	0
Gender	13
Married	3
Dependents	13
Education	0
Self_Employed	28
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	13
Credit_History	44
Property_Area	0
Loan_Status	0
dtype:	int64

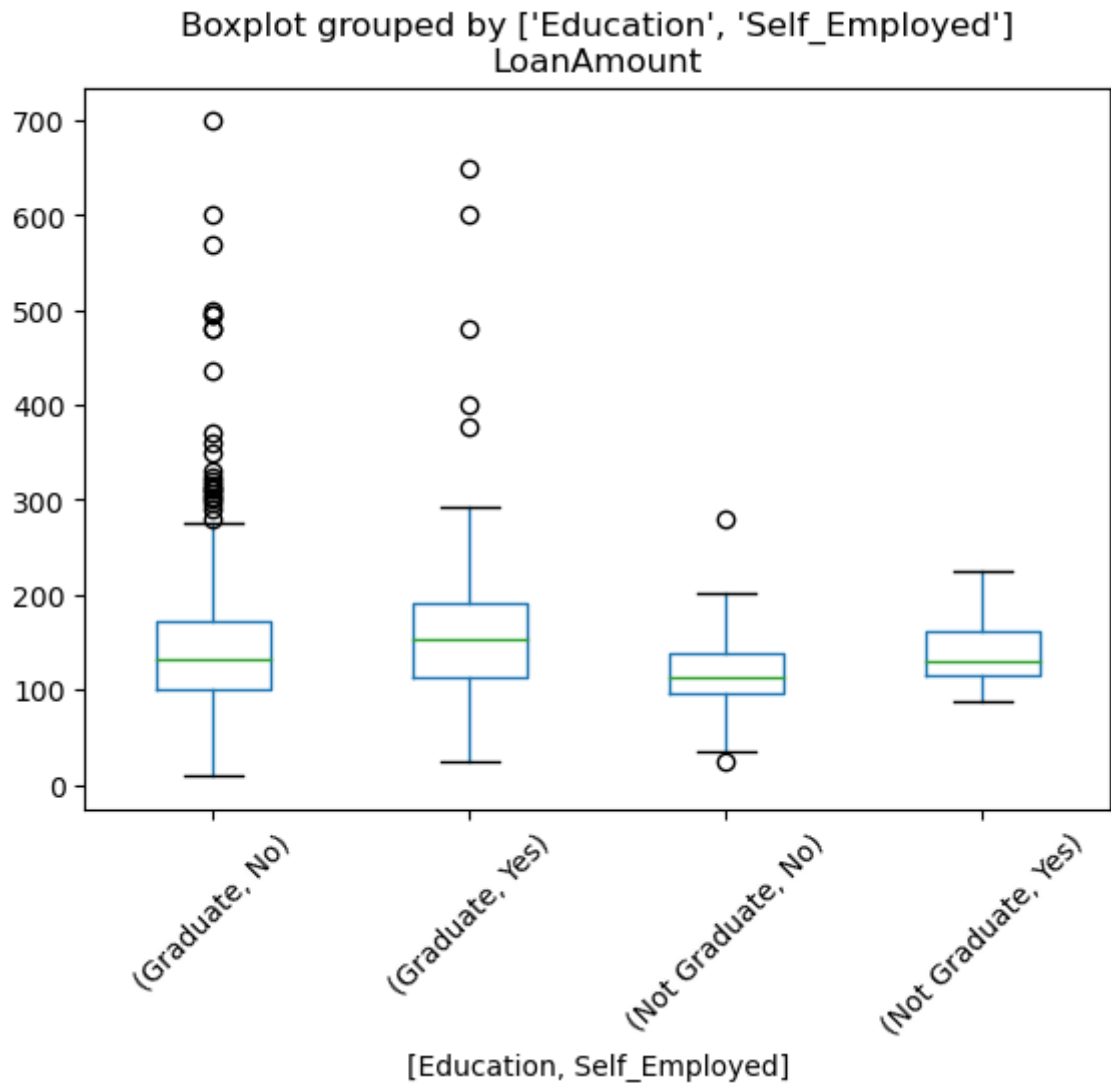
In [38]: `data.shape`

Out[38]: (550, 13)

In [39]: `data.to_csv('new_train.csv')`

In [40]: `data.boxplot(column='LoanAmount', by = ['Education', 'Self_Employed'], grid=False, rot = 45, fontsize = 10)`

Out[40]: <AxesSubplot:title={'center':'LoanAmount'}, xlabel=['Education, Self\_Employed']>



```
In [41]: data['Self_Employed'].value_counts()
```

```
Out[41]: No      445  
        Yes       77  
        Name: Self_Employed, dtype: int64
```

```
In [42]: data['Self_Employed'].fillna('No', inplace=True)
```

```
In [43]: data['Self_Employed'].value_counts()
```

```
Out[43]: No      473  
        Yes       77  
        Name: Self_Employed, dtype: int64
```

```
In [44]: data.apply(lambda x: sum(x.isnull()), axis=0)
```

```
Out[44]: Loan_ID      0
Gender      13
Married      3
Dependents   13
Education    0
Self_Employed  0
ApplicantIncome  0
CoapplicantIncome  0
LoanAmount    0
Loan_Amount_Term  13
Credit_History  44
Property_Area    0
Loan_Status      0
dtype: int64
```


```
In [45]: data.describe()
```

```
Out[45]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
<b>count</b>	550.000000	550.000000	550.000000	537.000000	506.000000
<b>mean</b>	5497.836364	1610.992582	148.527307	341.787709	0.839921
<b>std</b>	6263.552850	2938.590680	86.649203	64.796521	0.367042
<b>min</b>	150.000000	0.000000	9.000000	12.000000	0.000000
<b>25%</b>	2843.000000	0.000000	101.250000	360.000000	1.000000
<b>50%</b>	3815.000000	1149.000000	130.000000	360.000000	1.000000
<b>75%</b>	5844.000000	2297.250000	167.750000	360.000000	1.000000
<b>max</b>	81000.000000	41667.000000	700.000000	480.000000	1.000000


## Treating outliers in the dataset

```
In [46]: plt.hist(data['LoanAmount'], 20, facecolor='b')
plt.xlabel('Loan Amount')
plt.ylabel('Count')
plt.title('Histogram for loan amount')
plt.grid(True)
plt.show()
```




```
In [47]: data.boxplot(column='LoanAmount')
```

```
Out[47]: <AxesSubplot:>
```




```
In [48]: data['LoanAmount_log'] = np.log(data['LoanAmount'])  
#data['LoanAmount_log'].hist(bins = 20)
```

```
In [49]: plt.hist(data['LoanAmount_log'], 20, facecolor='b')
plt.xlabel('Loan Amount')
plt.ylabel('Count')
plt.title('Histogram for loan amount')
plt.grid(True)
plt.show()
```



```
In [50]: data.boxplot(column='LoanAmount_log')
```

```
Out[50]: <AxesSubplot:>
```



In [51]: `data.head()`

Out[51]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	

In [52]: `data.describe()`

Out[52]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	550.000000	550.000000	550.000000	537.000000	506.000000	550
mean	5497.836364	1610.992582	148.527307	341.787709	0.839921	0.538182
std	6263.552850	2938.590680	86.649203	64.796521	0.367042	0.500000
min	150.000000	0.000000	9.000000	12.000000	0.000000	0.000000
25%	2843.000000	0.000000	101.250000	360.000000	1.000000	0.000000
50%	3815.000000	1149.000000	130.000000	360.000000	1.000000	0.000000
75%	5844.000000	2297.250000	167.750000	360.000000	1.000000	0.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000	0.000000

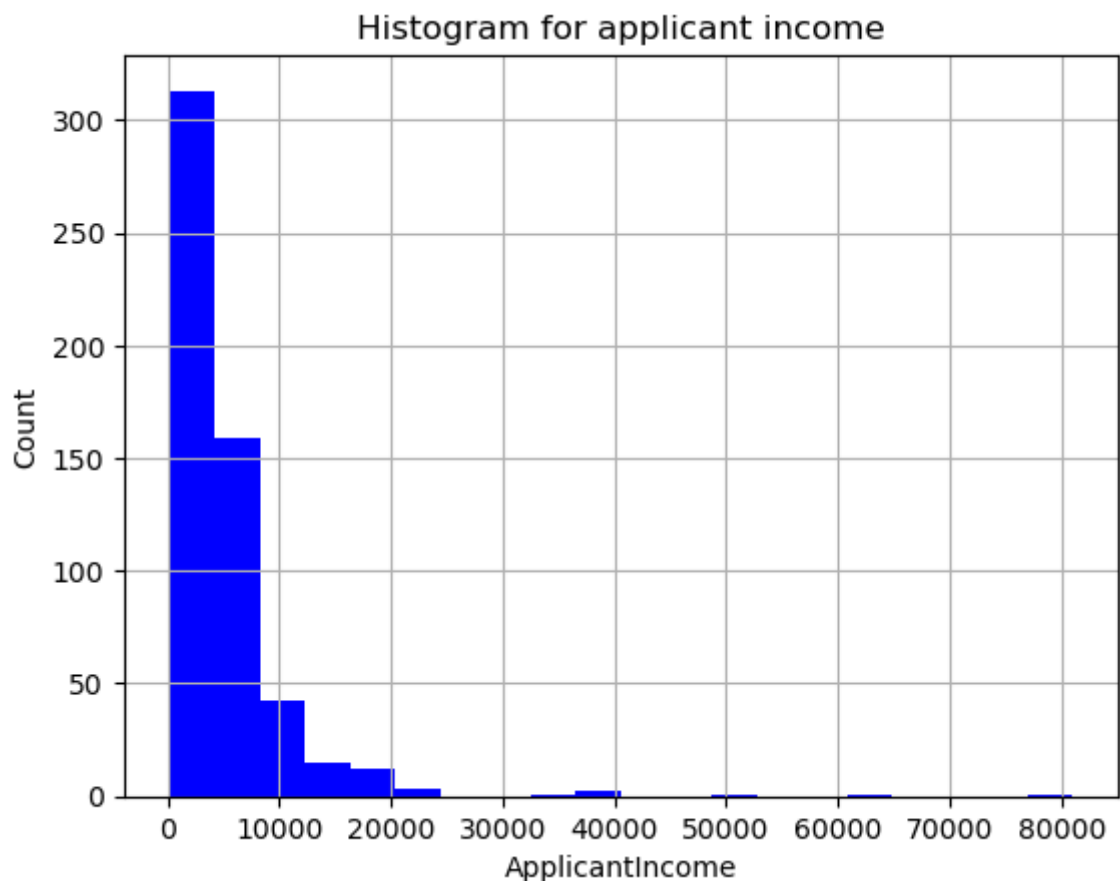
In [53]: `data = data.drop(['LoanAmount'], axis=1)`



# Try it yourself


I am checking outliers in "Loan\_Amount\_Term" and treating it with log transformation.

```
In [54]: plt.hist(data['ApplicantIncome'], 20, facecolor='b')  
plt.xlabel('ApplicantIncome')  
plt.ylabel('Count')  
plt.title('Histogram for applicant income')  
plt.grid(True)  
plt.show()
```




```
In [55]: data.boxplot(column='ApplicantIncome')
```

```
Out[55]: <AxesSubplot:>
```




```
In [56]: data['ApplicantIncome_log'] = np.log(data['ApplicantIncome'])
```

```
In [57]: plt.hist(data['ApplicantIncome_log'], 20, facecolor='b')  
plt.xlabel('ApplicantIncome')  
plt.ylabel('Count')  
plt.title('Histogram for applicant income')  
plt.grid(True)  
plt.show()
```



In [58]: `data.boxplot(column='ApplicantIncome_log')`

Out[58]: `<AxesSubplot:>`



In [59]: `data.head(10)`

Out[59]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	
5	LP002387	Male	Yes	0	Graduate	No	2425	
6	LP001038	Male	Yes	0	Not Graduate	No	4887	
7	LP002792	Male	Yes	1	Graduate	No	5468	
8	LP002478	NaN	Yes	0	Graduate	Yes	2083	
9	LP001137	Female	No	0	Graduate	No	3410	

In [60]: `data.describe()`

Out[60]:

	ApplicantIncome	CoapplicantIncome	Loan_Amount_Term	Credit_History	LoanAmount_log
<b>count</b>	550.000000	550.000000	537.000000	506.000000	550.000000
<b>mean</b>	5497.836364	1610.992582	341.787709	0.839921	4.872758
<b>std</b>	6263.552850	2938.590680	64.796521	0.367042	0.503418
<b>min</b>	150.000000	0.000000	12.000000	0.000000	2.197225
<b>25%</b>	2843.000000	0.000000	360.000000	1.000000	4.617584
<b>50%</b>	3815.000000	1149.000000	360.000000	1.000000	4.867534
<b>75%</b>	5844.000000	2297.250000	360.000000	1.000000	5.122471
<b>max</b>	81000.000000	41667.000000	480.000000	1.000000	6.551080

In [61]: `data['TotalIncome']=(data['ApplicantIncome']+data['CoapplicantIncome'])`  
`data.head()`

Out[61]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
<b>0</b>	LP001894	Male	Yes	0	Graduate	No	2620	
<b>1</b>	LP001814	Male	Yes	2	Graduate	No	9703	
<b>2</b>	LP002197	Male	Yes	2	Graduate	No	5185	
<b>3</b>	LP002739	Male	Yes	0	Not Graduate	No	2917	
<b>4</b>	LP001931	Female	No	0	Graduate	No	4124	

In [62]: `data['Gender'].fillna(data['Gender'].mode()[0], inplace = True)`  
*#0:gets the mode of each column, 1: for each row*  
`data['Married'].fillna(data['Married'].mode()[0], inplace = True)`  
`data['Dependents'].fillna(data['Dependents'].mode()[0], inplace = True)`  
`data['Loan_Amount_Term'].fillna(data['Loan_Amount_Term'].mode()[0], inplace = True)`  
`data['Credit_History'].fillna(data['Credit_History'].mode()[0], inplace = True)`

In [63]: `data.apply(lambda x: sum(x.isnull()), axis=0)`

Out[63]:

Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
Loan_Amount_Term	0
Credit_History	0
Property_Area	0
Loan_Status	0
LoanAmount_log	0
ApplicantIncome_log	0
TotalIncome	0
dtype:	int64

In [64]: `data.head()`

Out[64]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	Male	Yes	0	Graduate	No	2620	
1	LP001814	Male	Yes	2	Graduate	No	9703	
2	LP002197	Male	Yes	2	Graduate	No	5185	
3	LP002739	Male	Yes	0	Not Graduate	No	2917	
4	LP001931	Female	No	0	Graduate	No	4124	

In [65]: `data.shape`

Out[65]: (550, 15)

## Q5. Use LabelEncoder, to convert categorical variables into numeric. Hint: You will first need to identify categorial values.

In [66]: `from sklearn.preprocessing import LabelEncoder`

In [67]: `columns = list(data)  
print(columns)`

```
['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status', 'LoanAmount_log', 'ApplicantIncome_log', 'TotalIncome']
```

In [68]: `data.dtypes`

```
Out[68]: Loan_ID          object
Gender          object
Married         object
Dependents      object
Education        object
Self_Employed   object
ApplicantIncome  int64
CoapplicantIncome float64
Loan_Amount_Term float64
Credit_History  float64
Property_Area    object
Loan_Status      object
LoanAmount_log   float64
ApplicantIncome_log float64
TotalIncome      float64
dtype: object
```

In [69]: `columns = list(data.select_dtypes(exclude=['float64', 'int64']))`

In [70]: `c_columns = ['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']`

In [71]: `le = LabelEncoder()  
for i in c_columns:`

```
#print(i)
data[i] = le.fit_transform(data[i])
```

In [72]: `data.head()`

Out[72]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	1	1	0	0	0	2620	
1	LP001814	1	1	2	0	0	9703	
2	LP002197	1	1	2	0	0	5185	
3	LP002739	1	1	0	1	0	2917	
4	LP001931	0	0	0	0	0	4124	

I excluded the 'Loan\_ID' from been encoded as it is a unique list of identifiers for each Loan

In [73]: `#from sklearn.preprocessing import StandardScaler`  
`from sklearn.preprocessing import normalize`

In [74]: `original_data = data.copy()`  
`original_data.head()`

Out[74]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	1	1	0	0	0	2620	
1	LP001814	1	1	2	0	0	9703	
2	LP002197	1	1	2	0	0	5185	
3	LP002739	1	1	0	1	0	2917	
4	LP001931	0	0	0	0	0	4124	

In [75]: `original_data[0:5]`

Out[75]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	1	1	0	0	0	2620	
1	LP001814	1	1	2	0	0	9703	
2	LP002197	1	1	2	0	0	5185	
3	LP002739	1	1	0	1	0	2917	
4	LP001931	0	0	0	0	0	4124	

In [76]: `data[0:5]`

Out[76]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplic
0	LP001894	1	1	0	0	0	2620	
1	LP001814	1	1	2	0	0	9703	
2	LP002197	1	1	2	0	0	5185	
3	LP002739	1	1	0	1	0	2917	
4	LP001931	0	0	0	0	0	4124	

In [77]: `data_for_norm = data.drop(['Loan_ID', 'Loan_Status'], axis=1)`

In [78]: `normalized_data = normalize( data_for_norm )`

In [79]: `print(normalized_data[0:5])`

```
[1.68095792e-04 1.68095792e-04 0.00000000e+00 0.00000000e+00
 0.00000000e+00 4.40410974e-01 3.73676945e-01 6.05144850e-02
 1.68095792e-04 1.68095792e-04 8.42266706e-04 1.32307014e-03
 8.14087919e-01]
[7.28499787e-05 7.28499787e-05 1.45699957e-04 0.00000000e+00
 0.00000000e+00 7.06863344e-01 0.00000000e+00 2.62259923e-02
 7.28499787e-05 1.45699957e-04 3.43742542e-04 6.68776675e-04
 7.06863344e-01]
[1.36211271e-04 1.36211271e-04 2.72422541e-04 0.00000000e+00
 0.00000000e+00 7.06255438e-01 0.00000000e+00 4.90360574e-02
 1.36211271e-04 1.36211271e-04 6.86971343e-04 1.16508652e-03
 7.06255438e-01]
[2.19006276e-04 2.19006276e-04 0.00000000e+00 2.19006276e-04
 0.00000000e+00 6.38841307e-01 1.17387364e-01 7.88422593e-02
 2.19006276e-04 0.00000000e+00 9.17560683e-04 1.74730017e-03
 7.56228671e-01]
[0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 7.05762571e-01 0.00000000e+00 6.16087598e-02
 1.71135444e-04 1.71135444e-04 8.12026066e-04 1.42463050e-03
 7.05762571e-01]
```

In [80]: `normalized_data.shape`

Out[80]: (550, 13)

In [81]: `data.shape`

Out[81]: (550, 15)

In [82]: `normalized_data = pd.DataFrame(normalized_data, columns=data_for_norm.columns)`

In [83]: `normalized_data.head()`

Out[83]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	0.000168	0.000168	0.000000	0.000000	0.0	0.440411	0.37367
1	0.000073	0.000073	0.000146	0.000000	0.0	0.706863	0.00000
2	0.000136	0.000136	0.000272	0.000000	0.0	0.706255	0.00000
3	0.000219	0.000219	0.000000	0.000219	0.0	0.638841	0.11738
4	0.000000	0.000000	0.000000	0.000000	0.0	0.705763	0.00000

In [84]: `normalized_data['Loan_ID'] = data['Loan_ID']`

In [85]: `normalized_data['Loan_Status'] = data['Loan_Status']`

In [86]: `normalized_data.head()`

Out[86]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	0.000168	0.000168	0.000000	0.000000	0.0	0.440411	0.37367
1	0.000073	0.000073	0.000146	0.000000	0.0	0.706863	0.00000
2	0.000136	0.000136	0.000272	0.000000	0.0	0.706255	0.00000
3	0.000219	0.000219	0.000000	0.000219	0.0	0.638841	0.11738
4	0.000000	0.000000	0.000000	0.000000	0.0	0.705763	0.00000

In [87]: `normalized_data.describe()`

Out[87]:


	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
count	550.000000	550.000000	550.000000	550.000000	550.000000	550.000000	550.000000
mean	0.000115	0.000090	0.000101	0.000035	0.000016	0.574214	0.149447
std	0.000079	0.000082	0.000160	0.000075	0.000046	0.149447	0.007024
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.007024	0.465190
25%	0.000059	0.000000	0.000000	0.000000	0.000000	0.619191	0.706149
50%	0.000120	0.000091	0.000000	0.000000	0.000000	0.707105	
75%	0.000168	0.000153	0.000175	0.000000	0.000000		
max	0.000483	0.000373	0.001148	0.000483	0.000323		

## For Fun

In [88]: `normalized_data['LoanAmount_log'].hist(bins=100)`

Out[88]: `<AxesSubplot:>`





```
In [89]: from sklearn.model_selection import train_test_split
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.tree import export_graphviz
from sklearn.metrics import ConfusionMatrixDisplay
#import pydotplus
```

```
In [90]: columns = list(normalized_data.columns)
columns
```

```
Out[90]: ['Gender',
'Married',
'Dependents',
'Education',
'Self_Employed',
'ApplicantIncome',
'CoapplicantIncome',
'Loan_Amount_Term',
'Credit_History',
'Property_Area',
'LoanAmount_log',
'ApplicantIncome_log',
'TotalIncome',
'Loan_ID',
'Loan_Status']
```

```
In [91]: normalized_data.head()
```

Out[91]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	0.000168	0.000168	0.000000	0.000000	0.0	0.440411	0.37367
1	0.000073	0.000073	0.000146	0.000000	0.0	0.706863	0.00000
2	0.000136	0.000136	0.000272	0.000000	0.0	0.706255	0.00000
3	0.000219	0.000219	0.000000	0.000219	0.0	0.638841	0.11738
4	0.000000	0.000000	0.000000	0.000000	0.0	0.705763	0.00000

In [92]:

```
features = normalized_data.drop(['Loan_ID', 'Loan_Status'], axis = 1)
classes = pd.DataFrame(normalized_data['Loan_Status'])
```

In [93]:

```
print('Features:')
print(features.head())

print('Classes:')
print(classes.head())
```

Features:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	\
0	0.000168	0.000168	0.000000	0.000000	0.0	0.440411	
1	0.000073	0.000073	0.000146	0.000000	0.0	0.706863	
2	0.000136	0.000136	0.000272	0.000000	0.0	0.706255	
3	0.000219	0.000219	0.000000	0.000219	0.0	0.638841	
4	0.000000	0.000000	0.000000	0.000000	0.0	0.705763	

	CoapplicantIncome	Loan_Amount_Term	Credit_History	Property_Area	\
0	0.373677	0.060514	0.000168	0.000168	
1	0.000000	0.026226	0.000073	0.000146	
2	0.000000	0.049036	0.000136	0.000136	
3	0.117387	0.078842	0.000219	0.000000	
4	0.000000	0.061609	0.000171	0.000171	

	LoanAmount_log	ApplicantIncome_log	TotalIncome
0	0.000842	0.001323	0.814088
1	0.000344	0.000669	0.706863
2	0.000687	0.001165	0.706255
3	0.000918	0.001747	0.756229
4	0.000812	0.001425	0.705763

Classes:


	Loan_Status
0	1
1	1
2	1
3	0
4	1

In [94]:

```
normalized_data.head(10)
```

Out[94]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	0.000168	0.000168	0.000000	0.000000	0.00000	0.440411	0.37367
1	0.000073	0.000073	0.000146	0.000000	0.00000	0.706863	0.00000
2	0.000136	0.000136	0.000272	0.000000	0.00000	0.706255	0.00000
3	0.000219	0.000219	0.000000	0.000219	0.00000	0.638841	0.11738
4	0.000000	0.000000	0.000000	0.000000	0.00000	0.705763	0.00000
5	0.000171	0.000171	0.000000	0.000000	0.00000	0.414720	0.40018
6	0.000144	0.000144	0.000000	0.000144	0.00000	0.706149	0.00000
7	0.000117	0.000117	0.000117	0.000000	0.00000	0.638480	0.12050
8	0.000130	0.000130	0.000000	0.000000	0.00013	0.270819	0.53084
9	0.000000	0.000000	0.000000	0.000000	0.00000	0.705143	0.00000



In [95]: `normalized_data.shape`

Out[95]: (550, 15)

In [96]: `normalized_data.shape`

Out[96]: (550, 15)

In [97]: `from matplotlib import pyplot`

In [98]: `x_train, x_test, y_train, y_test = train_test_split(features, classes, test_size=`  
`random_state = 46)`

`print(x_train.shape, x_test.shape)`

(368, 13) (182, 13)


In [99]: `decisionTree = DecisionTreeClassifier(criterion='entropy')`  
`print(decisionTree)`

DecisionTreeClassifier(criterion='entropy')

In [100...]: `dtc_model = decisionTree.fit(x_train, y_train)`

In [101...]: `# feature importance`  
`importance = dtc_model.feature_importances_`  
`for i,v in enumerate(importance):`  
`print('Feature: %0d, Score: %.5f' % (i,v))`  
`# Barchat for feature importance`  
`pyplot.bar([x for x in range(len(importance))], importance)`  
`pyplot.show()`

Feature: 0, Score: 0.09070  
 Feature: 1, Score: 0.08610  
 Feature: 2, Score: 0.01503  
 Feature: 3, Score: 0.01095  
 Feature: 4, Score: 0.05621  
 Feature: 5, Score: 0.04769  
 Feature: 6, Score: 0.02707  
 Feature: 7, Score: 0.07054  
 Feature: 8, Score: 0.32666  
 Feature: 9, Score: 0.04490  
 Feature: 10, Score: 0.06708  
 Feature: 11, Score: 0.10648  
 Feature: 12, Score: 0.05058




```
In [102... prediction = dtc_model.predict(x_test)
```

```
In [103... y_true = le.inverse_transform(y_test["Loan_Status"])
y_pred = le.inverse_transform(prediction)
```

```
In [104... cm = confusion_matrix(y_true, y_pred)
labels = ['N', 'Y']
ConfusionMatrixDisplay(cm, display_labels=labels).plot()
```

```
Out[104]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2a44da195b0>
```



```
In [105... print(classification_report(y_true, y_pred))
```

	precision	recall	f1-score	support
N	0.44	0.60	0.51	48
Y	0.84	0.72	0.78	134
accuracy			0.69	182
macro avg	0.64	0.66	0.64	182
weighted avg	0.73	0.69	0.71	182

```
In [106... graphviz_path = 'C:/Program Files/Graphviz/bin/'
```


```
In [107... import os
os.environ["PATH"] += os.pathsep + graphviz_path
```

```
In [108... from graphviz import Source
from sklearn import tree
graph = Source( tree.export_graphviz(dtc_model, out_file=None, feature_names=featu
```

```
In [109... from cairosvg import svg2png
from IPython.display import Image

svg2png(bytestring=graph.pipe(format='svg'),write_to='output.png')
Image("output.png")
```

Out[109]:



## Report:

Q6. Based on the feature importance, select a different set of features to build another decision tree model. You should aim to improve the result of the baseline model.

The features I am selecting are: Credit\_History, LoanAmount\_log, TotalIncome

In [110...]

```
columns = list(normalized_data.columns)
columns
```

```
Out[110]: ['Gender',
'Married',
'Dependents',
'Education',
'Self_Employed',
'ApplicantIncome',
'CoapplicantIncome',
'Loan_Amount_Term',
'Credit_History',
'Property_Area',
'LoanAmount_log',
'ApplicantIncome_log',
'TotalIncome',
'Loan_ID',
'Loan_Status']
```

```
In [111...] features = normalized_data.drop(['ApplicantIncome', 'Loan_ID', 'Loan_Status', 'Gender'])
classes = pd.DataFrame(normalized_data['Loan_Status'])
```

```
In [112...] print('Features:')
print(features.head())
print('Classes:')
print(classes.head())
```

```
Features:
   Credit_History  LoanAmount_log  TotalIncome
0         0.000168         0.000842         0.814088
1         0.000073         0.000344         0.706863
2         0.000136         0.000687         0.706255
3         0.000219         0.000918         0.756229
4         0.000171         0.000812         0.705763
Classes:
   Loan_Status
0            1
1            1
2            1
3            0
4            1
```

```
In [113...] from matplotlib import pyplot
```

```
In [114...] x_train, x_test, y_train, y_test = train_test_split(features, classes, test_size=
                                                    random_state = 46)
print(x_train.shape, x_test.shape)
(368, 3) (182, 3)
```

```
In [115...] decisionTree = DecisionTreeClassifier(criterion='entropy')
print(decisionTree)
```

```
DecisionTreeClassifier(criterion='entropy')
```


```
In [116...] dtc_model2 = decisionTree.fit(x_train, y_train)
```

```
In [117...] # feature importance
importance = dtc_model2.feature_importances_
for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))
# Barchat for feature importance
pyplot.bar([x for x in range(len(importance))], importance)
pyplot.show()
```

Feature: 0, Score: 0.44795

Feature: 1, Score: 0.21494

Feature: 2, Score: 0.33711




```
In [118... prediction = dtc_model2.predict(x_test)
```

```
In [119... y_true = le.inverse_transform(y_test["Loan_Status"])  
y_pred = le.inverse_transform(prediction)
```

```
In [120... cm = confusion_matrix(y_true, y_pred)  
labels = ['N', 'Y']  
ConfusionMatrixDisplay(cm, display_labels=labels).plot()
```

```
Out[120]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2a44df2f3d0>
```





```
In [121... print(classification_report(y_true, y_pred))
```

	precision	recall	f1-score	support
N	0.48	0.65	0.55	48
Y	0.85	0.75	0.80	134
accuracy			0.72	182
macro avg	0.67	0.70	0.67	182
weighted avg	0.76	0.72	0.73	182

```
In [122... svg2png(bytestring=graph.pipe(format='svg'),write_to='output2.png')
Image("output2.png")
```

The diagram illustrates a decision tree for credit risk assessment. The root node is 'Credit\_History <= 0.0' (entropy = 0.927, samples = 368). The tree branches based on various features, including 'Married', 'TotalIncome', 'Property\_Area', 'LoanAmount\_log', 'Gender', 'Loan\_Amount\_Term', 'Self\_Employed', 'Education', 'Coefficientincome', 'Dependents', and 'Loan\_Amount\_Term'. Each internal node displays the split condition, the resulting entropy, the number of samples in each child node, and the predicted value. The tree structure is complex, with many nodes leading to leaf nodes that provide the final predicted values and sample counts.

In the first model (dtc\_model), all eleven features in the dataset were considered. However, in the second model (dtc\_model2), I selected three features that I deemed crucial in making informed decisions when granting loans to applicants. These features include Credit History, Loan Amount, and Total Income. It is vital to determine total income to ensure they are capable of repaying the loan. Furthermore, the loan amount should also be taken into consideration to ensure that it is within their means to repay based on their income. Lastly, Credit History provides a record of how an applicant has managed their credit in the past, including their total debt and payment timeliness. The bar chart of importance in In[124] also reiterates that Credit History is the most important followed by the Total income.

localhost:8888/nbconvert/html/OneDrive - University of Wolverhampton/Data Mining and Informatics/Week 2/Workshop 2 - Predicting loan case... 34/36

approved) are lower than those of class 'Y' (approved). This means that the model has difficulty in identifying the negative class, and it tends to misclassify a significant number of negative examples as positive. The macro-average F1-score of 0.65 indicates that the model's performance is suboptimal in terms of class imbalance, and there is room for improvement.

The second classification report shows an overall accuracy of 70%, which is the same as the previous report. However, the precision, recall, and F1-score for class 'N' have improved significantly, and the model now correctly identifies a higher percentage of negative examples. The precision, recall, and F1-score for class 'Y' are still high, indicating that the model is good at identifying the positive examples. The macro-average F1-score has also improved to 0.66, indicating a better balance between the two classes.

In conclusion, the second model with the reduced feature set performs better than the first model. However, the overall performance of both models is still suboptimal, and there is room for improvement in terms of identifying the negative class.

## Q8. Discuss the result based on the evaluation matrix (max 250 words).

The two confusion matrices from Out[104] and Out[128] represents the performance of a machine learning model in the prediction of well the model is handling its prediction. The rows in the matrix correspond to the actual class labels, while the columns correspond to the predicted class labels.

In the first confusion matrix, we can see that the first model correctly predicted the negative class(Not granted loans) (N) 29 times, but incorrectly predicted it 19 times. Similarly, the model correctly predicted the positive class(Granted loans) (Y) 98 times, but incorrectly predicted it 36 times. Overall, this confusion matrix suggests that the model is better at predicting the positive class than the negative class, as evidenced by the higher number of True Positives and False Negatives.

In the second confusion matrix, we can see that the second model correctly predicted the negative class(not granted) 31 times, but incorrectly predicted it 17 times. Similarly, the model correctly predicted the positive class(Granted loans) 97 times, but incorrectly predicted it 37 times. Compared to the first confusion matrix, this matrix suggests that the second model is slightly better at predicting the negative class than the first model, as evidenced by the higher number of True Negatives and lower number of False Negatives.

Overall, based on the two confusion matrices, we can conclude that the second machine learning model performed slightly better than the first one in terms of predicting the negative class.

## References

Syed Nor, S.H., Ismail, S. and Yap, B.W., 2019. Personal bankruptcy prediction using decision tree model. *Journal of Economics, Finance and Administrative Science*, 24(47), pp.157-170.

Cho, S., Hong, H. and Ha, B.C., 2010. A hybrid approach based on the combination of variable selection using decision trees and case-based reasoning using the Mahalanobis distance: For bankruptcy prediction. *Expert Systems with Applications*, 37(4), pp.3482-3488.